

1. General Remarks

This assignment is an introduction to MPI programming. You will be using the ece mpi cluster composed of 6 nodes:

`en-openmpiXX.ece.cornell.edu`

where `XX = 02, 03, 04, 05, 06, 07` Each node is a Ubuntu single socket server with 16 cores. Servers are connected by a 10GB router.

You will submit your jobs to SLURM scheduler from the head node `en-openmpi00.ece.cornell.edu`.

You should be able to login to any of the nodes using your Cornell netid. When you login your directory will be visible from any of the nodes so there is no need to copy files among the nodes.

IMPORTANT: The cluster is for ece5720 work only. You are not allowed to run any jobs on the cluster which are not ece5720 jobs.

2. Compile and execute

SURM requires two files `your_jobname.sub` and `your_jobname.sh`. Examples for these files are given below and also can be found on Canvas.

`your_jobname.sub`:

```
#!/bin/bash
#SBATCH -J mpi_hello          # Job name
#SBATCH -o output/mpi_hello.o%j # stdout output file(%j expands to jobId)
#SBATCH -e output/mpi_hello.e%j # stderr output file(%j expands to jobId)
#SBATCH --nodes=6               # number of nodes requested
#SBATCH --ntasks=12              # number of tasks to be configured for.
#SBATCH --tasks-per-node=2       # number of tasks to run on each node.
#SBATCH --cpus-per-task=1        # number of cpus needed by each task
#SBATCH --get-user-env           # retrieve the users login environment.
#SBATCH -t 00:10:00              # Run time (hh:mm:ss)
#SBATCH --mem-per-cpu=1000        # memory required per allocated CPU
#SBATCH --partition=six           # Which queue it should run on.
```

`/home/your_main_directory/mpi/mpi_hello.sh`

You need to create a subdirectory where your output errors (if any) will be written to. In the case above it is the `output` subdirectory.

`nodes` is the number of servers that you want to use. In our case the maximal number of servers is 6.

`ntasks` is the size returned by `MPI_Comm_size` in `MPI_COMM_WORLD`

`tasks-per-node` times `nodes` should equal `ntasks`

`cpus_per_task` usually we want only one MPI process per core.

`partition` at the moment there are 3 partitions:

- two nodes `en-openmpi02` and `en-openmpi03` which are identical
- gpu nodes `en-openmpi04` to `en-openmpi07` which are identical and equipped with GPUs
- six nodes `en-openmpi02` to `en-openmpi07` which is a default partition (that is if you do not specify a partition the system will assume `six` partition)

You also need to create `your_jobname.sh`: file

```
#!/bin/bash
mpirun -np 12 ./mpi_hello --mca opal_warn_on_missing_libcuda 0
```

Note that `-np xx` (in the example given `xx = 12`) must be equal to `ntasks`.

The flag `--mca opal_warn_on_missing_libcuda 0` discards error related to CUDA code, for this assignment we do not need CUDA support.

Before you submit a job for execution you have to compile it:

```
mpicc mpi_hello.c -o mpi_hello
```

plus all flags that are needed for compilation.

The next step is to submit the job

```
sbatch mpi_hello.sub
```

The system will assign a number to your job. You can find the output in the directory `output` where you identify the output by the number given to the job.

3. Warm up.

You should be aware that there are two levels of communication among different cores,

- within a single socket (any node),
- between different nodes.

To get an idea how fast the communication is, one can send progressively longer messages from one process to another, send the message back and time this so called "ping-pong" exchange.

MPI provides its own way for measuring time by means of `MPI_Wtime`

```
double start_time, end_time;
start_time = MPI_Wtime();
/* code to be timed */
end_time = MPI_Wtime();
```

You can find the resolution of the clock by the following directive

```
double tick;
tick = MPI_Wtick();
```

Split processes into two equal size groups and form pairs,

```
if (rank < numtasks/2)
    dest = src = numtasks/2 + rank;
if (rank >= numtasks/2)
    dest = src = rank - numtasks/2;
```

For the first group

```
start_time = MPI_Wtime();
MPI_Isend(&buffer, n, MPI_CHAR, dest, tag, MPI_COMM_WORLD, &reqs[0]);
MPI_Irecv(&buffer, n, MPI_CHAR, src, tag, MPI_COMM_WORLD, &reqs[1]);
MPI_Waitall(2, reqs, stats);
end_time = MPI_Wtime();
/* calculate average bandwidth for different length messages */
```

For the second group execute

```
MPI_Irecv(&msgbuf, n, MPI_CHAR, src, tag, MPI_COMM_WORLD, &reqs[1]);
MPI_Isend(&msgbuf, n, MPI_CHAR, dest, tag, MPI_COMM_WORLD, &reqs[0]);
MPI_Waitall(2, reqs, stats);
```

Send messages of length $n = 10^k$ with $k = 0, 2, 3, 4, 5, 6$. Measure the exchange time and avarage by the length of the message. This will give you an average time per per byte t_w ($1/t_w$ is the bandwidth).

If you send a single character this will give you some sense of latency (or startup time t_s). Clearly the "ping-pong" should be repeated several times (say 20) and the total time needs to be averaged.

Remember that the results may differ depending on which partition you use and how many tasks are requested.

Present your timing results on a graph and in a table, and describe your observations.

Repeat the experiment using blocking send and receive.

Your timing code must be saved in a file `your_netid_timing.c`.

4. One sided Jacobi for computing the EVD

For graph partitioning we need an eigensolver. The problem can be stated as follows.

For a given matrix A we want to find orthogonal U and V , and diagonal Σ such that

$$AV = U\Sigma \quad (1)$$

The equation (1)i, after transposition, becomes

$$V^T A^T = \Sigma^T U^T \quad (2)$$

We prefer (2) as in that form Givens rotations are applied to rows. And in "C" arrays are stored by rows.

We approximate the decomposition (2) by performing so called sweeps. A sweep is a product of $\frac{n(n-1)}{n}$ Givens rotations. A sweep in turn is a product of compound rotations. And a compound rotation is a product of $\frac{n}{2}$ independent rotations. These are rotations that operate on pairs of consecutive rows. A general form of a compound rotations for $n = 8$ is the following

$$J = \begin{pmatrix} c_0 & s_0 & & & & & & \\ -s_0 & c_0 & & & & & & \\ & & c_1 & s_1 & & & & \\ & & -s_1 & c_1 & & & & \\ & & & & c_2 & s_2 & & \\ & & & & -s_2 & c_2 & & \\ & & & & & & c_3 & s_3 \\ & & & & & & -s_3 & c_3 \end{pmatrix}$$

The parameters s_i and c_i are such that $s_i^2 + c_i^2 = 1$ (thuss, these are sines and cosines of a common angle).

Say that after k sweeps we have matrices $V^{(k)}$ and $A^{(k)}$. We want to get $V^{(k+1)}$ and $A^{(k+1)}$. For that we compute another sweep composed of compound rotations.

For a compound rotation we select pairs $(2p, 2p+1)$ of rows $a_{2p}^{(k)}$ and $a_{2p+1}^{(k)}$, $p = 0, \dots, \frac{n}{2} - 2$, and form a 2×2 matrix

$$B = \begin{pmatrix} b_{2p,2p} & b_{2p,2p+1} \\ b_{2p+1,2p} & b_{2p+1,2p+1} \end{pmatrix}$$

where

$$b_{2p,2p} = a_{2p}^{(k)}(a_{2p}^{(k)})^T, \quad b_{2p,2p+1} = a_{2p}^{(k)}(a_{2p+1}^{(k)})^T = b_{2p+1,p}, \quad b_{2p+,2p+1} = a_{2p+1}^{(k)}(a_{2p+1}^{(k)})^T$$

Now we find cosine and sine c_p and s_p so B is diagonalized,

$$\begin{pmatrix} c_p & s_p \\ -s_p & c_p \end{pmatrix}^T \begin{pmatrix} b_{2p,2p} & b_{2p,2p+1} \\ b_{2p+1,2p} & b_{2p+1,2p+1} \end{pmatrix} \underbrace{\begin{pmatrix} c_p & s_p \\ -s_p & c_p \end{pmatrix}}_{G_p} = \begin{pmatrix} \hat{b}_{2p,2p} & 0 \\ 0 & \hat{b}_{2p+1,2p+1} \end{pmatrix}$$

Now pairs of rows of $V^{(k)}$ and $A^{(k)}$ are updated by multiplying them by Givens matrices G_p^T . This completes a compound rotation.

The next thing we want to do is to permute the working matrix so another set of row pairs is created. This is done by the Brent-Luk scheme. In the scheme, indices are shifted along a ring as illustrated in Figure 1 (and in the notes for Lecture 14).




Figure 1.

Shifts are executed by the following pseudo code

```

index = 1;

// Shift all elements in bottom row to the left
for (i = 0; i < n/2 - 1; i++) {
    A[index] = A[index + 2];
    index += 2;
}

// Shift the top right to bottom right
A[index] = A[index - 1];
index -= 1;

// Shift top row from left to right
for (i = 0; i < n/2 - 1; i++) {
    A[index] = A[index - 2];
    index -= 2;
}

```

```

// Fill in the temp value
A[0] = col_temp;

// Make sure that the element A[1] moves to A[2]

```

After the shifts are executed another compound rotation is generated. This is repeated $n - 1$ times (as there are $\frac{N(n-1)}{2}$ rotations in a sweep).

Sweeps are repeated until a (user defined) maximum number of sweeps are executed, or when desired accuracy is achieved.

There are number of ways to check for accuracy. One way is to compute

$$\text{error} = \text{off}(A^{(k)}(A^{(k)})^T)$$

where off was defined in Lecture 14. If error is smaller than a desired threshold, the iterations will stop. However the computation of off requires matrix-matrix multiplication and is expensive. Rather than performing matrix-matrix multiplication for the entire matrix, one can cheat a bit by performing matrix-matrix multiplication on submatrices that reside in all PEs. The local off can be accumulated by `MPI_Allreduce` on all nodes and compared to the threshold. Also, we will start computing off after $s = \min(8, \log_2(n))$ sweeps have been performed.

5 Format.

We assume a ring topology for P MPI processes and a block row distribution of the data matrix A .

Your algorithm should accept the following arguments:

- `n` - dimension of the matrix
- `max_sweeps` - maximal number of sweeps allowed
- `tol` - measure of orthogonality of $A^{(k)}$.

Populate the $n \times n$ matrix A with random numbers using `drand48` random numbers generator.

- Your code must be well documented so any person with limited knowledge of C could understand what the code is doing.
- For a fixed number of PEs you need to run your code for progressively larger matrices.
- For a fixed size of matrix A you need to run your code for progressively larger number of PEs.
- Your timing data must include at least the following cases: $n = 128, 512, 1024, P = 8, 32, 64$.

- For each case record speed-up (or slow down) over a sequential algorithm ($P = 1$ and no MPI). If there is a slow down, identify the parts of the code that contribute to the slow down.
- Your Jacobi code must be saved in the file `your_netid_jacobi.c`. You also need to submit `your_netid_jacobi.sub` and `your_netid_jacobi.sh` scripts.
- Write a document that describes how your programs work.
- Sketch the key elements of your internode communication strategy and describe how synchronization is achieved.
- Your findings, a discussion of results, graphs and tables should be saved in a file `your_netid_writeup.pdf` (NO *.docx files please).
- Please present your tables and graphs in a way so they are easily readable. Use the scientific format (`x.xxxxeyy` which show 3 most significant decimal digits and 2 digits of exponent).
- The codes, the write-up and files `.sh` and `.sub` (see Part 2) must be archived as `your_netid.zip` or `your_netid.tar` files.
- NOTE: When we unzip your submission we want to see only the files
 - `your_netid_writeup.pdf`,
 - `your_netid_timings.c`,
 - `your_netid_jacobi.c`,
 - `your_netid_jacobi.sh`.
 - `your_netid_jacobi.sub`.
- Please NO other files or subdirectories.
- If you rely on resources outside lecture notes but publicly available, please cite sources in your write-up.
- You can discuss the homework with your classmates but the final work must be your own (no code sharing).